

Kafka Connect

Raghavi, Janaswamy

Sr.Principal Engineer (TLCP)

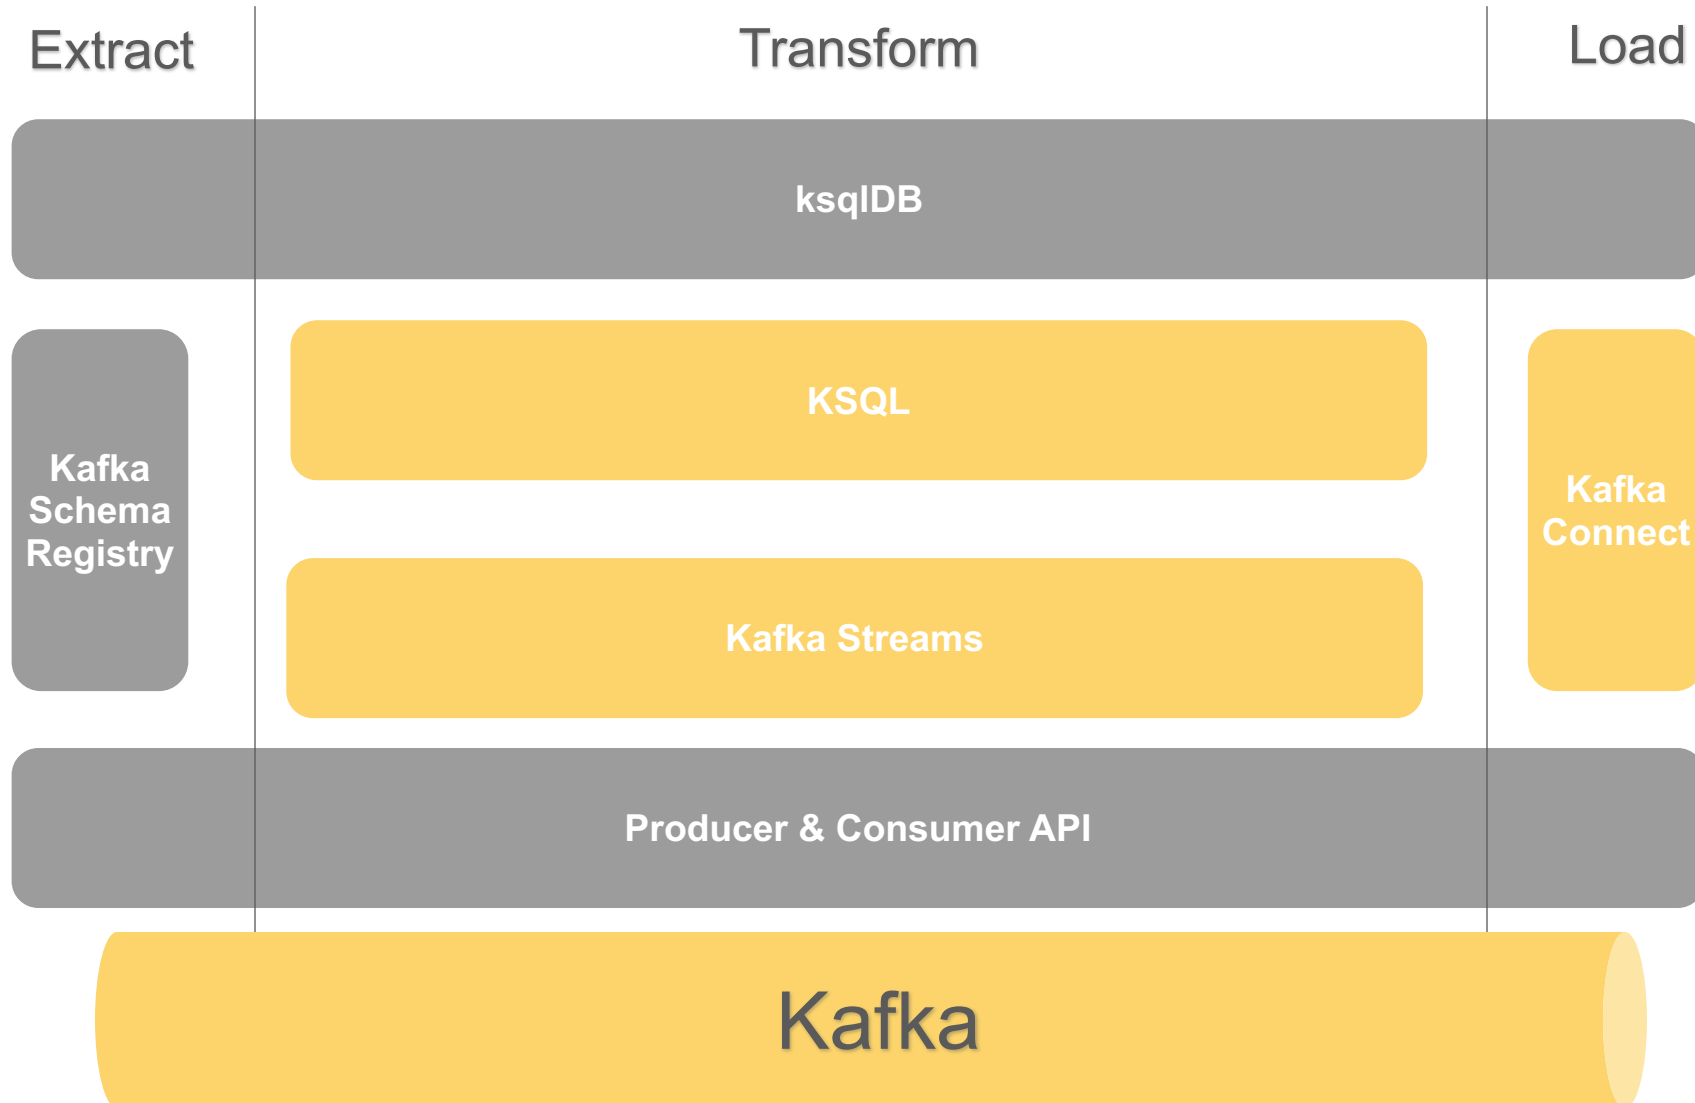
raghavi.janaswamy@optum.com



Agenda

- Kafka & Kafka connect Concepts
- Lab 1: Kafka platform Setup - Local Docker Environment
- Lab 2: Kafka Connect – Build Source Connector
- Lab 3: Kafka Connect – Build Sink Connector
- Kafka Connect implementations in OptumCare & P360
- Pointers to the Advanced Concepts
- Q&A

Kafka Ecosystem

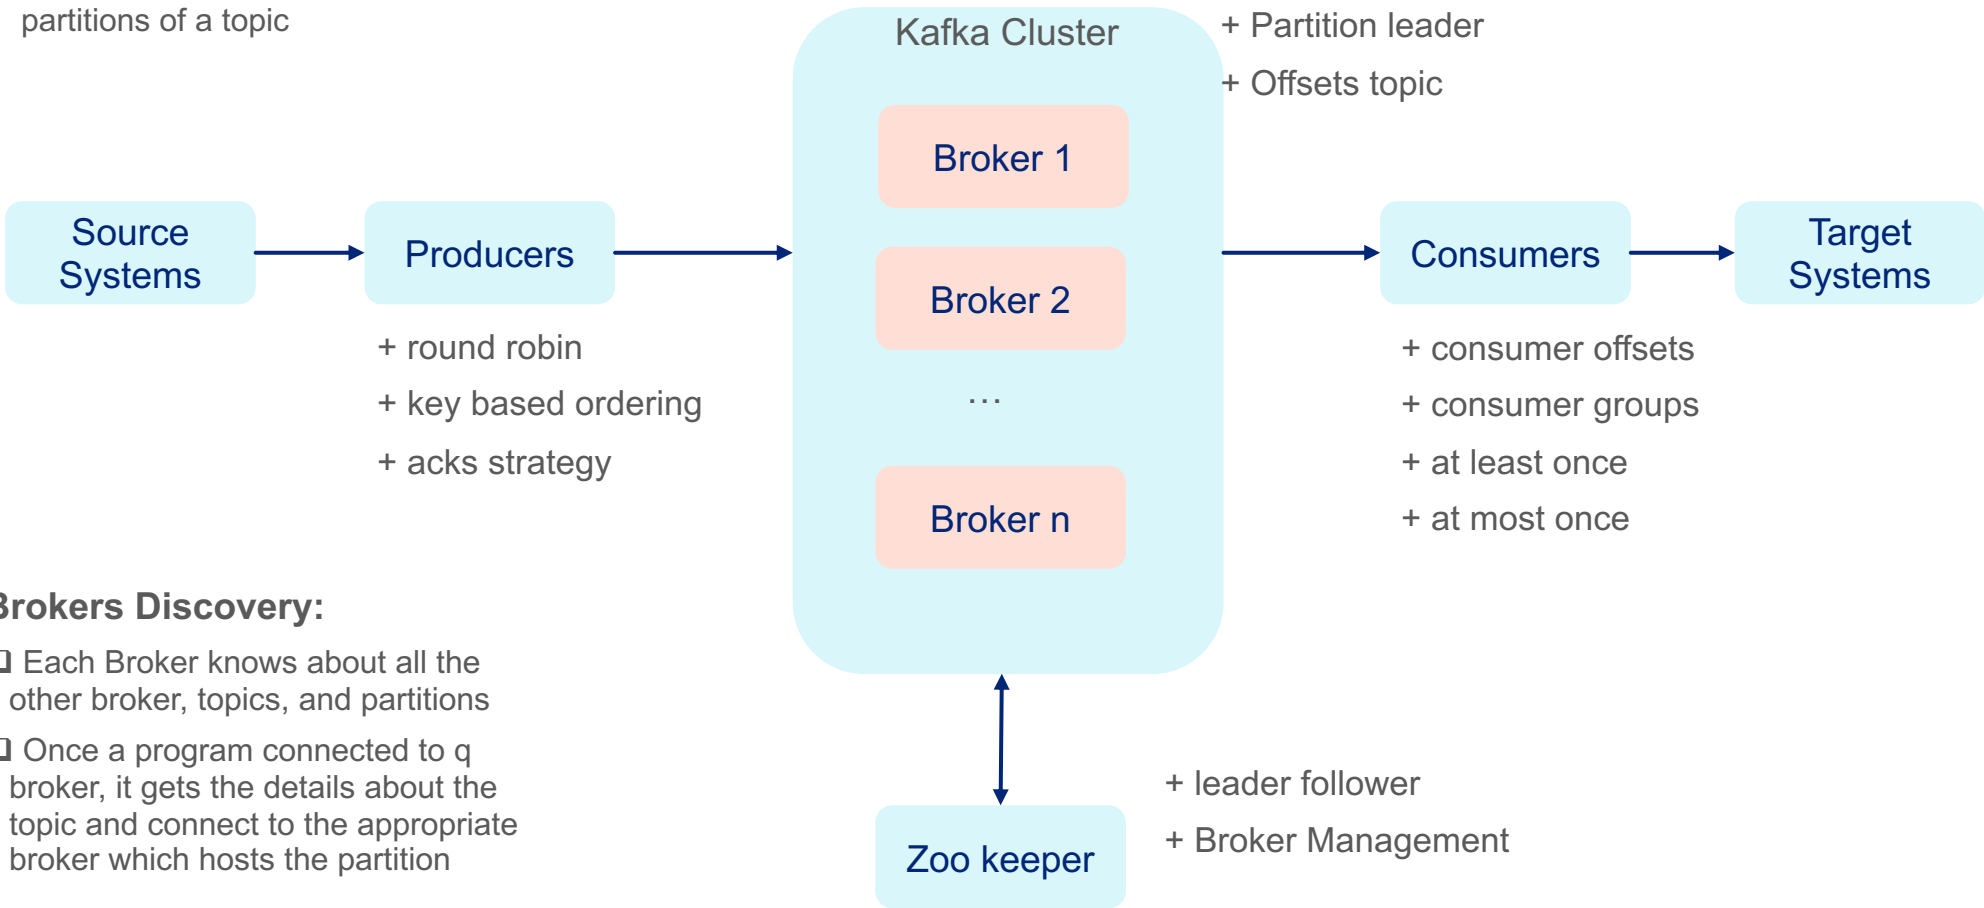


Kafka Core - Terminology

Brokers and Topics:

- Basically a server in a Kafka cluster.
- Each broker will contain certain partitions of a topic

- + topics
- + partitions
- + replication
- + Partition leader
- + Offsets topic



Brokers Discovery:

- Each Broker knows about all the other broker, topics, and partitions
- Once a program connected to q broker, it gets the details about the topic and connect to the appropriate broker which hosts the partition

Kafka - Internals

☐ Topic Replication:

- ☐ Distributed world needs replication.
- ☐ At any point in time, there will be one broker for a partition which will serve as a leader.
- ☐ Leader sends and receives data. Other brokers just stay in sync.

☐ Producer - Data Acknowledgement

- ☐ Producer can choose to receive acknowledgment from the Broker
 - ☐ Producer won't wait for acknowledgment
 - ☐ Producer will wait for acknowledgment
 - ☐ Producer will wait for acknowledgment for leader and replicas

☐ Message keys

- ☐ Data is sent in round robin across partitions
- ☐ Data for a particular key is always sent to a particular partition. Key hashing.

☐ Consumers and Consumer Groups:

- ☐ Consumers can read from multiple partitions for a single topic in parallel.
- ☐ Consumer read data in groups, called consumer groups
- ☐ Two consumers of the same group will read from mutually exclusive partitions


☐ Consumer Offsets/ Acknowledgement

- ☐ Stored to a Kafka topic "__consumer_offsets"
- ☐ Consumers can continue from where they left

☐ Delivery Semantics

- ☐ Consumer choose when to commit offsets
 - ☐ Commit once received
 - ☐ Commit once received and processed

Terminology - Kafka Topics and Messages



☐ Topic:

- ☐ A particular stream of data.
- ☐ Every record is called a message.
- ☐ A topic will have a name
- ☐ Data is kept for a limited duration

☐ Offset:

- ☐ Unique id given to a message with in a partition
- ☐ First message will have an offset of "0"
- ☐ Offset will NEVER go back to zero, even after the earlier messages are deleted.

☐ Partitions:


- ☐ Topics are split in partitions.
- ☐ Partitions are ordered
- ☐ Data once written to the partition, it CANNOT be changed. It's immutable.

☐ Message

- ☐ To identify a message - use the partition id & offset together
- ☐ Message in a particular partition is chronologically ordered.
- ☐ Messages across partitions in a topic are NOT chronologically ordered

What is Kafka Connect

Free, open-source component of Kafka ecosystem that provides extensible interface to connect to different data sources & load to or from Kafka in scalable and reliable way.



Advantages

- **Standalone & Distributed Compute:** Kafka connect can run as a cluster providing multiple workers on which a job can be run.
- **Extensible:** Provides extensible interface that can be leveraged to create different source and sink connectors. Many free/open-source sink connectors can be plugged into a connect cluster for connecting to different sources.
- **Configurable:** All jobs can be configured at the start of the cluster or can be provided via REST API to a running cluster instance.
- **Reliable:** Provides at least once processing & retry capabilities with logging and metrics exposed for monitoring & alerting.
- **Data Transformations:** Provides capability to perform light weight transformations on the data.

How Kafka Connect works

Terminologies

Worker: Processes running on instances that make a Kafka cluster.

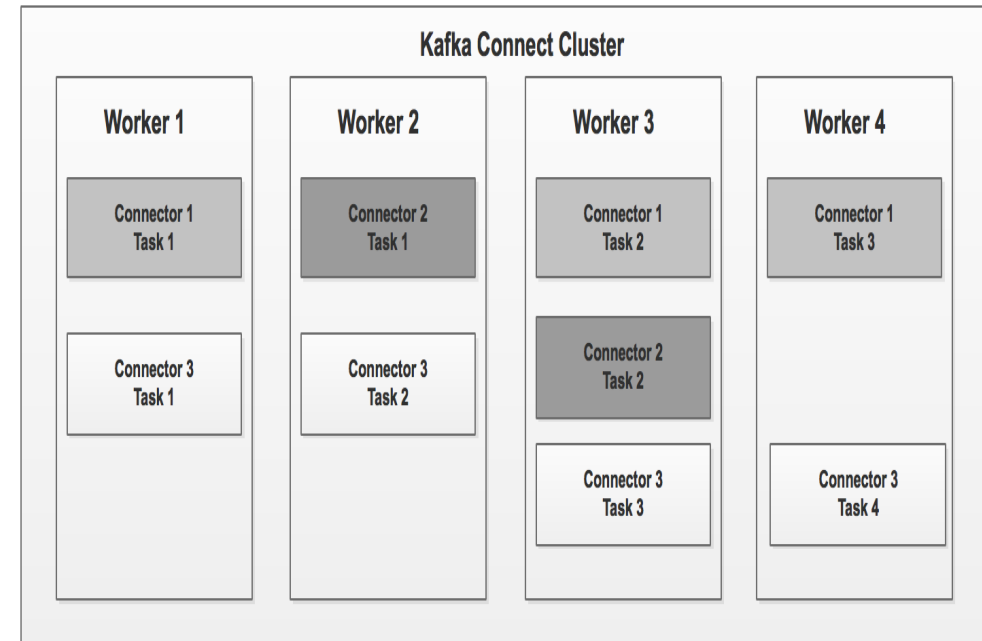
■ Types:

■ Standalone

- Single process running on an instance.
- Scalability/Fault-tolerance not achieved.
- Very useful in testing/small extractions.

■ Distributed

- Scalable solution to have parallel processing.
- Embedded fault tolerance based on inherent Kafka Infrastructure utilization.

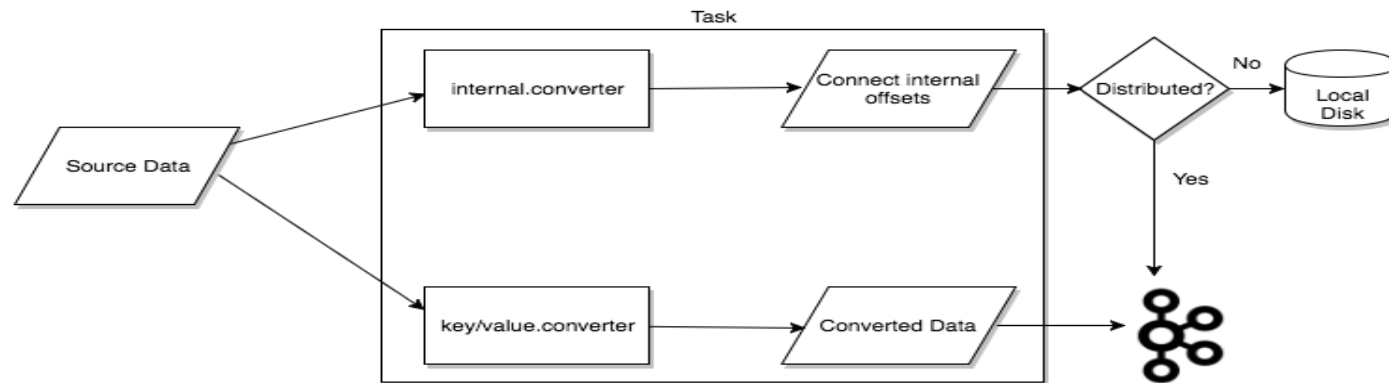


Connect Internals: Tasks

- **Tasks** : Implementation that actually copy the data

Connector Jobs -> Task 1, Task 2, Task n

- **Internal Topics (Compacted)** : Task State storage
 - **Config Topic**: Stores all connector configuration.
 - **Offset Topic**: Stores all source connector offsets to support incremental data extraction. Please note that Sink connector works on default consumer_offsets topic for consumer offset storage.
 - **Status Topic**: Stores all connector statuses.



Type of Kafka Connectors

Connectors: Job configuration that connects Kafka to different sources/destinations

- **Source Connectors :** Imports data into Kafka Topic

Ex: JDBC Source

JMS Source

- **Sink Connectors :** Exports data from Kafka Topic

JDBC Sink

Elasticsearch Sink

Amazon S3 Sink

Sample Source Connector

```
{
  "name": "mysql-curl-claimdop-volumes",
  "config": {
    "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",
    "incrementing.column.name": "Id",
    "transforms.createKey.type": "org.apache.kafka.connect.transforms.ValueToKey",
    "connection.password": "optum",
    "tasks.max": "2",
    "transforms": "createKey,extractInt",
    "transforms.extractInt.type": "org.apache.kafka.connect.transforms.ExtractField$Key",
    "batch.max.rows": "10",
    "table.types": "TABLE",
    "table.poll.interval.ms": "2000",
    "table.whitelist": "claims_dop_metrics_volumes_sink_test",
    "mode": "incrementing",
    "topic.prefix": "kaas.pdp.dev.p360-",
    "transforms.extractInt.field": "tin",
    "connection.user": "root",
    "transforms.createKey.fields": "tin",
    "poll.interval.ms": "2000",
    "name": "mysql-curl-claimdop-volumes",
    "connection.url": "jdbc:mysql://apsrt4811.uhc.com:30898/optum",
    "client.id": "kaas.pdp.dev"
  }
}
```

Sample C* Sink Connector

```
{
  "name": "cassandra-sink-claims-dop-volumes",
  "config": {
    "connector.class": "com.datamountaineer.streamreactor.connect.cassandra.sink.CassandraSinkConnector",
    "tasks.max": "1",
    "topics": "kaas.pdp.p360-claims-dop-metrics-volumes-volumes-output",
    "connect.cassandra.contact.points": "apvrd22444.uhc.com,apvrd22443.uhc.com,apvrd22445.uhc.com",
    "connect.cassandra.port": "9042",
    "connect.cassandra.username": "pdp",
    "connect.cassandra.password": "<password>",
    "connect.cassandra.key.space": "claims_streams",
    "connect.cassandra.kcql": "INSERT INTO claim_tin_dop_metrics SELECT tin AS provider_tin,ownedOrAffiliated AS owned_affiliated,claimTypeCode AS claim_type_code,lineOfBusiness AS line_of_business,submissionType AS submission_type,claimNetworkType AS claim_in_out_network_type,claimSystemId AS claim_system_id,fundingTypeCode AS funding_type_code,dateOfProcess AS date_of_process,hcoName AS hco_name,claimsApproved AS approved_count,claimsDenied AS denied_count FROM kaas.pdp.p360-claims-dop-metrics-volumes-volumes-output;",
    "value.converter.schema.registry.url": "http://kaas-test-schema-registry-a.optum.com",
    "value.converter": "io.confluent.connect.avro.AvroConverter",
    "key.converter.schema.registry.url": "http://kaas-test-schema-registry-a.optum.com",
    "key.converter": "io.confluent.connect.avro.AvroConverter",
    "connect.cassandra.default.value": "UNSET",
    "connect.cassandra.ssl.enabled": "true",
    "connect.cassandra.trust.store.password": "<truststore_password>",
    "connect.cassandra.trust.store.path": "/mnt/pdp-cassandra-certs/cassandra.truststore.jks",
    "connect.cassandra.key.store.password": "<keystore_password>",
    "connect.cassandra.key.store.path": "/mnt/pdp-cassandra-certs/cassandra.keystore.jks",
    "connect.cassandra.ssl.client.cert.auth": true
  }
}
```

Kafka Connect: Converters

Purpose: Translate the data between source/sink systems and Kafka connector


- **Converter**: The format conversion while writing to/reading from Kafka.
 - AVROConverter, JSONConverter, StringConverter
- **Transform**: Simple/lightweight conversions to be applied on each individual message.
 - Cast, Drop, Flatten, InsertField, ValueToKey
- **Plugin**: Implementation classes which define the execution behavior of Connectors.

Kafka connect - REST APIs

- GET <http://localhost:8085/connectors>
- GET <http://localhost:8085/connector-plugins>
- POST <http://localhost:8085/connectors>
- GET http://localhost:8085/connectors/<connector_name>/status
- PUT http://localhost:8085/connectors/<connector_name>/<pause_resume_restart>
- DELETE http://localhost:8085/connectors/<connector_name>

Lab – Scenarios

- Lab 1 – Setup
Internal topics setup
docker-compose to setup kafka environment/database for source and sink
- Lab 2 – Kafka Connect Source
Kafka Source Connector setup
Verify data in the kafka topic
- Lab 3 – Kafka Connect Sink
Kafka Sink Connector Setup
Verify data in the Cassandra Database



Kafka Connect @ P360 / OptumCare

- OptumCare
- <https://github.optum.com/OCNP/OCNP-caredata-kconnect-source-facets>
- <https://github.optum.com/OCNP/OCNP-caredata-kconnect-sink-facets>
- Open source Kafka connector:
 - <https://github.optum.com/EnterpriseProviderPlatform/stream-reactor/tree/master/kafka-connect-cassandra>
- Deployment Image Creation:
 - <https://github.optum.com/EnterpriseProviderPlatform/pdp-kconnect-rdbms-cassandra>
- Example:
 - <https://github.optum.com/EnterpriseProviderPlatform/network-measures-kconnect-s3-sink-lab/tree/master/network-measures-kconnect-s3-sink-lab-chart>

Transformations (SMTs)

Purpose: Single Message transformations on individual messages.

| Cast: cast to a specific type

- "transforms": "Cast", "transforms.Cast.type": "org.apache.kafka.connect.transforms.Cast\$Value", "transforms.Cast.spec": "ID:string,score:float64"

ExtractField: Extract a specified field when not null

- "transforms": "extractField", "transforms.ExtractField.type": "org.apache.kafka.connect.transforms.ExtractField\$Key", "transforms.ExtractField.field": "id"

ExtractTopic: Replace topic with a new topic derived from its key/value

- "transforms": "keyFieldExample", "transforms.KeyFieldExample.type": "io.confluent.connect.transforms.ExtractTopic\$Value", "transforms.KeyFieldExample.field": "f3", "transforms.KeyFieldExample.skip.missing.or.null": true

Flatten: Flatten a nested data structure

- "transforms": "flatten", "transforms.flatten.type": "org.apache.kafka.connect.transforms.Flatten\$Value", "transforms.flatten.delimiter": "."

ReplaceField: Filter/replace fields

- "transforms": "ReplaceField", "transforms.ReplaceField.type": "org.apache.kafka.connect.transforms.ReplaceField\$Value", "transforms.ReplaceField.blacklist": "c2"

TimestampConverter:

- "transforms": "keyFieldExample", "transforms.KeyFieldExample.type": "io.confluent.connect.transforms.ExtractTopic\$Value", "transforms.KeyFieldExample.field": "f3", "transforms.KeyFieldExample.skip.missing.or.null": true

ValueToKey:

- "transforms": "ValueToKey", "transforms.ValueToKey.type": "org.apache.kafka.connect.transforms.ValueToKey", "transforms.ValueToKey.fields": "userId,city,state"

Error Handling:

How: Create Dead letter queue with below parameters

```
"errors.tolerance": "all",  
"errors.deadletterqueue.topic.name": "dlq_topic_name",  
"errors.deadletterqueue.topic.replication.factor": 1
```

```
errors.deadletterqueue.context.headers.enable = true
```

Q/A

Thank you

Feedback Link

<https://www.metricsthatmatter.com/url/u.aspx?74E258F91202202062>